

Brian_McKean_Stress_Management

March 8, 2017

1 Brian McKean

2 =====

3 BioTrak Health

<https://www.biotrakhealth.com/> ## Session Management ## Identify ineffective sessions during user trials # ===== ## Galvanize Data Science Immersion ## Capstone Project

First data exploration

```
In [1]: from docx import Document
        from docx.shared import Inches
        import pandas as pd
        import numpy as np
        import matplotlib.pyplot as plt
        from mpl_toolkits.axes_grid.parasite_axes import SubplotHost
        from mpl_toolkits.axes_grid1 import host_subplot
        import mpl_toolkits.axisartist as AA
        import seaborn as sns
        import requests
        import scipy.stats as stats
        from sklearn.ensemble import RandomForestClassifier
        from sklearn.model_selection import train_test_split
        from sklearn.metrics import precision_recall_curve
        from sklearn.externals import joblib
        from sklearn.model_selection import cross_val_score
        from sklearn.model_selection import cross_val_score
        from sklearn.metrics import precision_recall_curve
        from sklearn.ensemble import AdaBoostClassifier
        from sklearn.ensemble import GradientBoostingClassifier
```

3.1 Get Data from files

```
In [2]: # Read in data
        df_users = pd.read_json('users_02_26_2017.txt')
        df_sessions = pd.read_json('sessions_02_26_2017.txt')
```

```

print df_users.columns
print df_sessions.columns

Index([u'YOB', u'_id', u'email', u'first_name', u'gender', u'last_name',
       u'userid'],
      dtype='object')
Index([u'Build#', u'Device_info_id', u'End_GMT_time', u'End_local_time',
       u'GMT_Date', u'GMT_date', u'Local_date', u'Session_data_header',
       u'Session_name', u'Session_name:', u'Session_time_interval',
       u'Session_type', u'Start_GMT_time', u'Start_local_time',
       u'Start_time_zone', u'User_info_id', u'Version#', u'_id',
       u'session_data', u'user_info_id', u'userid'],
      dtype='object')

```

3.1.1 Initial Data Exploration

```

In [3]: print len(df_users), " users are set up"
        print len(df_sessions), "sessions are uploaded"

```

```

66  users are set up
356 sessions are uploaded

```

```

In [4]: # Read Ratings File
        df_ratings = pd.read_csv('data/ratings.csv')
        print "Number of ratings is ",len(df_ratings)

```

```

Number of ratings is  356

```

```

In [5]: print "Rated Bad:", df_ratings[df_ratings['rating']=='Bad'].count()

```

```

Rated Bad: Unnamed: 0      237
session_num      237
rating           237
dtype: int64

```

```

In [6]: print "Rated Good:",df_ratings[df_ratings['rating']=='Good'].count()

```

```

Rated Good: Unnamed: 0      119
session_num      119
rating           119
dtype: int64

```

```

In [7]: df_sessions = df_sessions.ix[:356,: ]

```

Set up scoring After a review session the project firmware engineer and I graded 356 sessions. There were 237 bad and 119 good ratings. a

In addition to coming up with the ratings we jointly decided on some heuristics for determining a good and bad sessions

1. If the session was less than 60 seconds the session is bad. It is possible that a user may have an properly operating device over the 60 seconds, but we consider that not an effective use as there is not enough time to use the sessions as designed.
 2. If the session, after 60 seconds, goes to 0 (Min) or 4095 (Max), the session is bad. Some of the sessions in this category show a significant amount of good data, but the fact that the user had problems idnicate someting that needs to be addressed.
- These two extremes indicate that an electrode has lost contact with the skin. Once a session has started the electrodes should continue to make proper contact until the end of the session.
3. Low signal values are not a problem.
 4. Signal may be smooth with gradual peaks or there may be sharp peaks. The peaks come from activities such as jaw clenching that may be part of a session or may simply be the user exploring the operation of the device.

Add New Features In order to incororate the heuristics I add the following features

1. min_len (T/F) – is the length at least 3600 units (60 seconds)
2. pegged_L (T/F) – at least one value of 0 after 60 seconds
3. pegged_H (T/F) – at least one value of 4095 afte 60 seconds

```
In [8]: # add ratings to session data
        df_sessions['rating'] = df_ratings['rating']
```

```
In [9]: print df_sessions.columns
```

```
Index([
            u'Build#',          u'Device_info_id',
            u'End_GMT_time',     u'End_local_time',
            u'GMT_Date',        u'GMT_date',
            u'Local_date',      u'Session_data_header',
            u'Session_name',     u'Session_name:',
            u'Session_time_interval', u'Session_type',
            u'Start_GMT_time',    u'Start_local_time',
            u'Start_time_zone',   u'User_info_id',
            u'Version#',         u'_id',
            u'session_data',     u'user_info_id',
            u'userid',          u'rating'],
      dtype='object')
```

A litte data exploration

```

In [10]: df_sessions.Session_name.unique()

Out[10]: array([u'Open', u'Breathing for Relief', u'Head and Neck Relax',
                u'3 Minute Relax', u'Deep Breathing', u'Forest Visualization',
                u'Deep Relax for Sleep', u'Elevator Excercise',
                u'Beach Visualization', u'Cloud Visualization',
                u'Extended Exercise', u'Breath and Relax', u'Progress Muscle Relax

In [11]: df_sessions.Session_type.unique()

Out[11]: array([u'open', u'managed'], dtype=object)

In [12]: df_groups = df_sessions.groupby('Session_name')

In [13]: df_groups['Device_info_id'].count()

Out[13]: Session_name
          3 Minute Relax      32
          Beach Visualization    9
          Breath and Relax      4
          Breathing for Relief   10
          Cloud Visualization    6
          Deep Breathing         6
          Deep Relax for Sleep   6
          Elevator Excercise     7
          Extended Exercise      1
          Forest Visualization   6
          Head and Neck Relax    10
          Open                  257
          Progress Muscle Relax  2
          Name: Device_info_id, dtype: int64

In [14]: df_results = df_sessions[['Session_name', 'userid']]

In [15]: # Score is True if session is bad
          # we are looking for bad sessions
          df_results['score'] = np.where(df_sessions['rating']=='Good', False, True)

/Users/brianmckean/anaconda2/lib/python2.7/site-packages/ipykernel/__main__.py:3: S
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/10min.html#copy-on-write
app.launch_new_instance()

```

We have to change format for the time series data. It starts as a string of comma delimited values. The last value is uploaded incorrectly in some sessions so we discard it.

```

In [16]: # Pull session data dictionary info into it's own clumns
          dfNew = df_sessions.join(pd.DataFrame(df_sessions["session_data"].to_dict

```

```
In [17]: dfNew.columns
```

```
Out[17]: Index([u'Build#', u'Device_info_id',
                u'End_GMT_time', u'End_local_time',
                u'GMT_Date', u'GMT_date',
                u'Local_date', u'Session_data_header',
                u'Session_name', u'Session_name:',
                u'Session_time_interval', u'Session_type',
                u'Start_GMT_time', u'Start_local_time',
                u'Start_time_zone', u'User_info_id',
                u'Version#', u'_id',
                u'session_data', u'user_info_id',
                u'userid', u'rating',
                u'average_data', u'average_data_count',
                u'baseline_data', u'data_interval',
                u'final_data', u'graphed_data',
                u'graphed_data_count', u'raw_data',
                u'raw_data_count'],
               dtype='object')
```

```
In [18]: # Start calculation of features
         # Convert session data from string to array of ints
         # -- Drop the last entry as there was a bug that put in a bad last entry
temp = dfNew['average_data'].str.replace(" ", "").str.split(',').str[:-1].m
```

```
In [19]: # Some of the time series are empty, if so insert an entry of one sample a
         # Signal is inverted -- need to change it
temp = temp.map(lambda x: np.array( [0] if not len(x) else x ))
temp = temp.map(lambda x: 4095 - x)
```

```
In [20]: # Get the mean
temp2 = temp.map(lambda x: x.mean())
```

```
In [21]: temp2.fillna(0,inplace=True)
df_results['avg_data_mean'] = temp2
```

```
/Users/brianmckean/anaconda2/lib/python2.7/site-packages/ipykernel/__main__.py:2: S
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/
from ipykernel import kernelapp as app
```

```
In [22]: # Get the max
t_max = temp.map(lambda x: x.max())
```

```
In [23]: df_results['avg_data_max'] = t_max
```

```
/Users/brianmckean/anaconda2/lib/python2.7/site-packages/ipykernel/__main__.py:1: S
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/
if __name__ == '__main__':
```

```
In [24]: # Get the min
         df_results['avg_data_min'] = temp.map(lambda x: x.min())
```

```
/Users/brianmckean/anaconda2/lib/python2.7/site-packages/ipykernel/__main__.py:2: S
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/
from ipykernel import kernelapp as app
```

```
In [25]: # Get the length
         df_results['avg_data_len'] = temp.map(lambda x: len(x))
```

```
/Users/brianmckean/anaconda2/lib/python2.7/site-packages/ipykernel/__main__.py:2: S
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/
from ipykernel import kernelapp as app
```

```
In [26]: # Get the standard deviation
         df_results['avg_data_std'] = temp.map(lambda x: x.std())
```

```
/Users/brianmckean/anaconda2/lib/python2.7/site-packages/ipykernel/__main__.py:2: S
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

```
See the caveats in the documentation: http://pandas.pydata.org/pandas-docs/stable/
from ipykernel import kernelapp as app
```

```
In [27]: df_results.columns
```

```
Out[27]: Index([ u'Session_name',          u'userid',          u'score', u'avg_data_me
              u'avg_data_max', u'avg_data_min', u'avg_data_len', u'avg_data_s
              dtype='object')
```

```
In [28]: # Load in session num
         df_results['session_num'] = df_results.index
```

```
/Users/brianmckean/anaconda2/lib/python2.7/site-packages/ipykernel/__main__.py:2: S
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/>
from ipykernel import kernelapp as app

```
In [29]: # Add heuristics
df_results['min_len'] = df_results['avg_data_len']>3600
df_results['pegged_low'] = [False if (len(x)<3600 or min(x[3600:])>0) else
df_results['pegged_high'] = [False if (len(x)<3600 or max(x[3600:])<4095)
```

```
/Users/brianmckean/anaconda2/lib/python2.7/site-packages/ipykernel/__main__.py:2: S
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/>
from ipykernel import kernelapp as app

```
/Users/brianmckean/anaconda2/lib/python2.7/site-packages/ipykernel/__main__.py:3: S
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/>
app.launch_new_instance()

```
/Users/brianmckean/anaconda2/lib/python2.7/site-packages/ipykernel/__main__.py:4: S
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/>

```
In [30]: df_results.columns
```

```
Out[30]: Index([ u'Session_name',          u'userid',          u'score', u'avg_data_me
            u'avg_data_max', u'avg_data_min', u'avg_data_len', u'avg_data_s
            u'session_num',          u'min_len',          u'pegged_low', u'pegged_hi
            dtype='object')
```

```
In [31]: df_key_columns = [x for x in list(df_results.columns) if x not in ['Sess
print df_key_columns
```

```
['avg_data_mean', 'avg_data_max', 'avg_data_min', 'avg_data_len', 'avg_data_std',
```

```
In [32]: df_res_save = df_results.copy()
```

```
In [33]: y = df_results['score'].as_matrix()
X_df = df_results[df_key_columns]
X = X_df.as_matrix()
```

```
In [34]: X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=0.33, random_state=42)
```

```
In [35]: df_results.describe()
```

```
Out[35]:
```

	avg_data_mean	avg_data_max	avg_data_min	avg_data_len	avg_data_
count	356.000000	356.000000	356.000000	356.000000	356.000
mean	2080.052353	3521.893258	1021.233146	13404.457865	602.631
std	997.930074	883.259504	1220.579913	34697.181848	424.079
min	0.000000	0.000000	0.000000	1.000000	0.000
25%	1317.220337	3357.250000	0.000000	1152.000000	264.431
50%	2090.857233	3901.000000	423.000000	3736.000000	508.647
75%	2837.760875	4047.000000	1945.250000	13222.000000	939.515
max	4095.000000	4095.000000	4095.000000	382019.000000	1667.963

	session_num
count	356.000000
mean	177.500000
std	102.912584
min	0.000000
25%	88.750000
50%	177.500000
75%	266.250000
max	355.000000

```
In [36]: clf = RandomForestClassifier(n_estimators=10, random_state=42).fit(X_train,
```

```
In [37]: scores = cross_val_score(clf, X_train, y_train, cv=10)
```

```
In [38]: print scores
        print scores.mean()
```

```
[ 0.83333333  0.79166667  0.95833333  0.91666667  0.79166667  0.875
  0.95833333  0.875          0.82608696  0.82608696]
0.865217391304
```

```
In [39]: pred = clf.predict(X_train)
```

```
In [40]: print clf.feature_importances_
```

```
for i in range(len(clf.feature_importances_)):
    print X_df.columns[i], clf.feature_importances_[i]
```

```
clf_save = clf
clf_name_save = X_df.columns
X_save = X_df.as_matrix()
```

```
[ 0.07482741  0.05206467  0.09356981  0.28436927  0.10351105  0.08949767
  0.15199783  0.14889914  0.00126317]
```



```
avg_data_mean 0.074827409527
avg_data_max 0.0520646674285
avg_data_min 0.0935698076556
avg_data_len 0.284369266812
avg_data_std 0.10351104636
session_num 0.0894976678294
min_len 0.151997825363
pegged_low 0.148899139956
pegged_high 0.00126316906828
```

```
In [41]: pred_test = clf.predict(X_test)
```

```
In [42]: clf.score(X_test,y_test)
```

```
Out[42]: 0.92372881355932202
```

```
In [43]: clf.get_params
```

```
Out[43]: <bound method RandomForestClassifier.get_params of RandomForestClassifier
max_depth=None, max_features='auto', max_leaf_nodes=None,
min_impurity_split=1e-07, min_samples_leaf=1,
min_samples_split=2, min_weight_fraction_leaf=0.0,
n_estimators=10, n_jobs=1, oob_score=False, random_state=42,
verbose=0, warm_start=False)>
```

```
In [44]: y_prob = clf.predict_proba(X_test)[:,-1:]
```

```
In [45]: missed = np.where([pred_test != y_test])[1]
```

```
In [46]: for idx in missed:
        print df_results.iloc[idx]['session_num']
```

```
10
20
29
56
64
67
71
88
115
```

```
In [47]: for i in range(len(missed)):
        print i, clf.predict_proba(X_test)[i]
```

```
0 [ 0.  1.]
1 [ 0.  1.]
2 [ 0.7  0.3]
```

```

3 [ 0.  1.]
4 [ 0.  1.]
5 [ 0.2  0.8]
6 [ 0.  1.]
7 [ 0.  1.]
8 [ 0.8  0.2]

```

```
In [ ]:
```

```
In [48]: precision, recall, thresholds = precision_recall_curve(y_test, y_prob)
```

```
In [49]: print precision, recall, thresholds
```

```

[ 0.75238095  0.84615385  0.90588235  0.94805195  0.97260274  1.          1.
  1.          1.          1.          1.          ] [ 1.          0.97468354  0.97468
  0.87341772  0.83544304  0.82278481  0.65822785  0.          ] [ 0.1  0.2  0.3  0.4

```

```
In [50]: df = pd.DataFrame()
```

```
In [51]: df['precision'] = precision
```

```

In [52]: print len(precision), len(recall), len(thresholds)
          thresholds = np.append(thresholds,1.0)
          print len(precision), len(recall), len(thresholds)

```

```
11 11 10
```

```
11 11 11
```

```

In [53]: df['recall'] = recall
          df['thresholds'] = thresholds

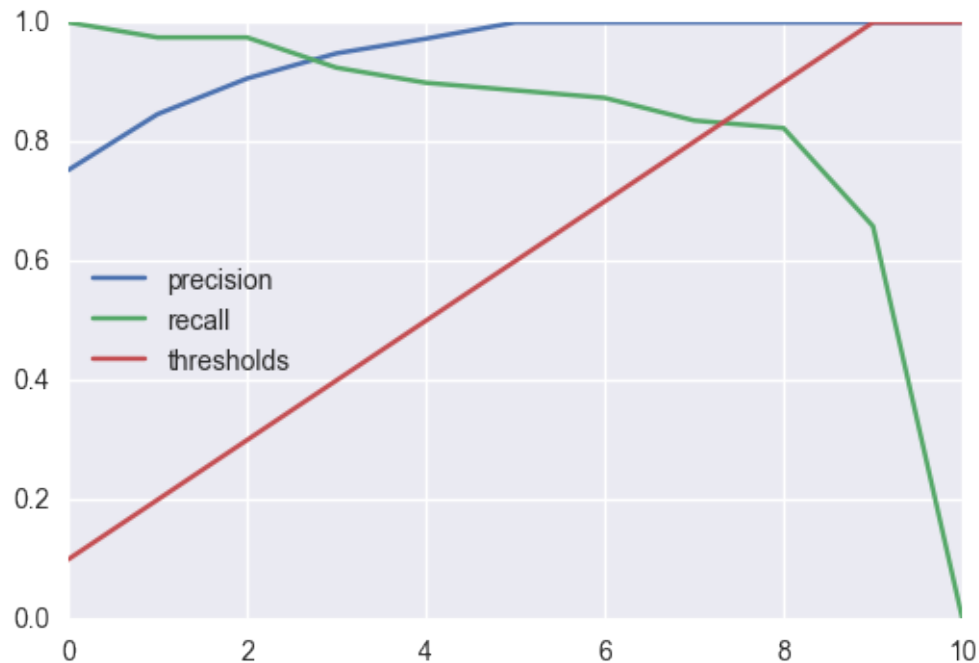
```

```

In [54]: %matplotlib inline
          df.plot()

```

```
Out[54]: <matplotlib.axes._subplots.AxesSubplot at 0x12e04b050>
```



```
In [55]: byScore = df_results.groupby('score').count()
```

```
In [56]: print byScore['avg_data_min']
```

```
score
False      119
True       237
Name: avg_data_min, dtype: int64
```

```
In [57]: ## Add in the cycles and largest swing features
```

```
def count_crosses(s, pts=10, intv=600):
    '''
    Looks at series
    for each 'intv' points
    - count crossings of 'pts' lookback moving average
    INPUTS:
        s = list of measurements
        pts = how many pts to collect for moving average
        intv = interval to measure crossing (600 = 10 sec)
    OUTPUTS
        crossings / intv = average crossing per interval
        biggest_move = on a crossing biggest move
```

```

'''
if len(s) < intv:
    return 0, 0
if pts >= intv:
    return 0,0
s = np.array(s)
crossing_counts = [0]
index = pts
max_swing = 0
up = True # True for last cross up, False for down
while (index + intv < len(s)):
    #print index
    interval_crossings = 0
    for i in range(intv):
        last_n = s[index-pts:index]
        avg = last_n.mean()
        if up and s[index] < avg:
            up = False
            interval_crossings += 1
            swing = abs(s[index]-s[index-1])
            if swing > max_swing:
                max_swing = swing
        if not up and s[index] > avg:
            up = True
            interval_crossings += 1
            swing = abs(s[index]-s[index-1])
            if swing > max_swing:
                max_swing = swing
        index += 1
    crossing_counts.append(interval_crossings)

time_sec = intv/60.0
return np.array(crossing_counts).mean()/(2*time_sec), max_swing

```

```

In [58]: df_results3 = df_res_save.copy()
cycles = []
swings = []

look_back = [1]
#look_back = [3]

for look in look_back:
    c = []
    s = []
    for idx, elem in enumerate(temp):

```

```

        a,b = count_crosses(elem, pts=look)
        c.append(a)
        s.append(b)
    cycles.append(c)
    swings.append(s)
    df_results3['cycles_{0}'.format(look)] = c
    df_results3['swings_{0}'.format(look)] = s

    #plt.hist(cycles)
    #plt.hist(swings)

```

```
In [59]: #df_results3
```


```
In [60]: df4 = df_results.copy()
         df4['cycles_1'] = df_results3['cycles_1']
         df4['swings_1'] = df_results3['swings_1']

```

```
In [ ]:
```

```
In [61]: sns_plot = sns.pairplot(df4)
         sns_plot.savefig("sns-pairplot.png")

```



```
In [62]: y2 = df4['score'].as_matrix()
X_df2 = df4[['avg_data_mean',
             'avg_data_max', 'avg_data_min', 'avg_data_std', 'avg_data_len',
             'cycles_1', 'swings_1', 'min_len', 'pegged_low', 'pegged_high']]
X2 = X_df2.as_matrix()

In [63]: X_train, X_test, y_train, y_test = train_test_split(
        X2, y2, test_size=0.33, random_state=42)

In [64]: clf = RandomForestClassifier(n_estimators=10, random_state=42).fit(X_train,
        y_train)
clf = clf.fit(X_train, y_train)

In [65]: clf_save = clf
clf_name_save = X_df2.columns
X_save = X_df2.as_matrix()
```

```
In [66]: scores = cross_val_score(clf, X_train, y_train, cv=10)

In [67]: print scores
          print np.average(scores)

[ 0.79166667  0.95833333  0.875          0.875          0.875          0.95833333
  0.91666667  0.95833333  0.82608696  0.82608696]
0.886050724638
```

```
In [68]: print clf.feature_importances_

          for i in range(len(clf.feature_importances_)):
              print X_df2.columns[i], clf.feature_importances_[i]

[ 0.07064774  0.05054319  0.09737334  0.10564576  0.25162631  0.1348587
  0.0572746   0.13476878  0.09468364  0.00257796]
avg_data_mean 0.070647739072
avg_data_max 0.050543185561
avg_data_min 0.0973733381049
avg_data_std 0.10564575551
avg_data_len 0.251626310731
cycles_1 0.134858702122
swings_1 0.0572745976741
min_len 0.134768777261
pegged_low 0.094683635819
pegged_high 0.00257795814473
```

```
In [ ]:
```

```
In [69]: precision, recall, thresholds = precision_recall_curve(y_test, y_prob)
          print precision, recall, thresholds

[ 0.75238095  0.84615385  0.90588235  0.94805195  0.97260274  1.          1.
  1.          1.          1.          1.          ] [ 1.          0.97468354  0.97468
  0.87341772  0.83544304  0.82278481  0.65822785  0.          ] [ 0.1  0.2  0.3  0.4
```


```
In [70]: fig = plt.figure()
          ax = fig.add_subplot(1,1,1)
          ax.set_ylabel('Precision and Recall',fontsize=20)
          x1 = np.arange(0,10)/5.0
          #ax.plot(x1,df['precision'])
          labels = [item.get_text() for item in ax.get_xticklabels()]
          labels[1] = 'Testing'
          ax.plot(df['precision'])
          ax.plot(df['recall'])
          #ax.plot(df['thresholds'])
```

```

ax.set_xlim(0,10)
#labels =
ax.set_xticklabels(x1)
ax.set_title('Precision and Recall Versus Probability Threshold',fontsize=
ax.set_xlabel('Probability Threshold',fontsize=20)
#plt.ylabel('Density',fontsize=18)
#plt.title('Average Frequency in Tension Measurement', fontsize=24)
ax.tick_params(axis='both',labelsize=16)
plt.legend(fontsize=20,loc=0)
plt.show()
plt.savefig('prec-recall.png',bbox_inches='tight')

```

Precision and Recall Versus Probability Threshold



<matplotlib.figure.Figure at 0x12819ba90>

```
In [71]: df4_t = df4[df4['score']==True]
```

```
In [72]: df4_f = df4[df4['score']==False]
```

```
In [73]: # Set up the matplotlib figure
```

```

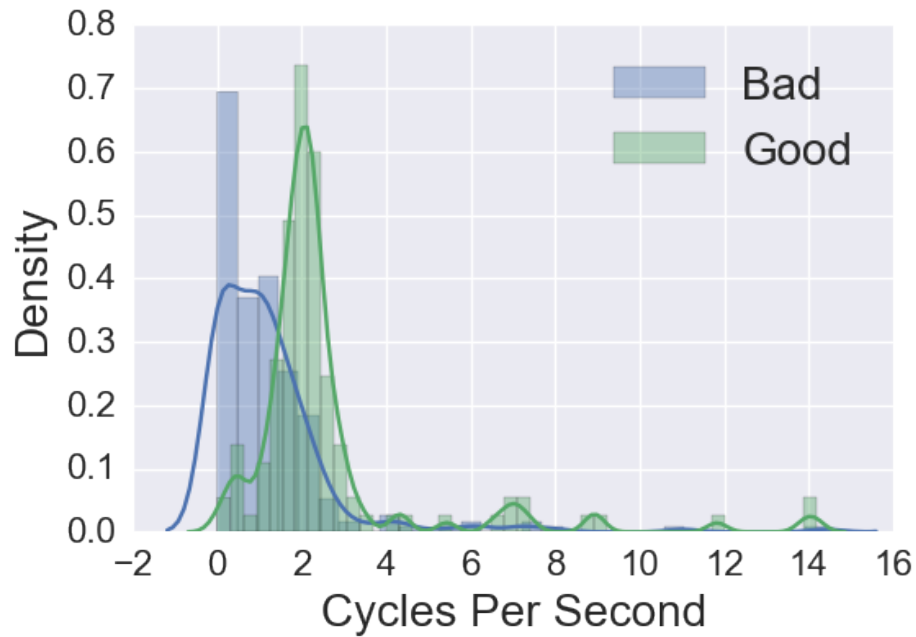
sns.distplot( df4_t['cycles_1'], label = 'Bad')
sns.distplot( df4_f['cycles_1'], label = 'Good')
plt.xlabel('Cycles Per Second',fontsize=20)
plt.ylabel('Density',fontsize=20)
plt.title('Average Frequency in Tension Measurement', fontsize=24, y=1.10)

```




```
plt.tick_params(axis='both',labelsize=16)
plt.legend(fontsize=20)
plt.savefig('feat-freq.png',bbox_inches='tight')
```

Average Frequency in Tension Measurement

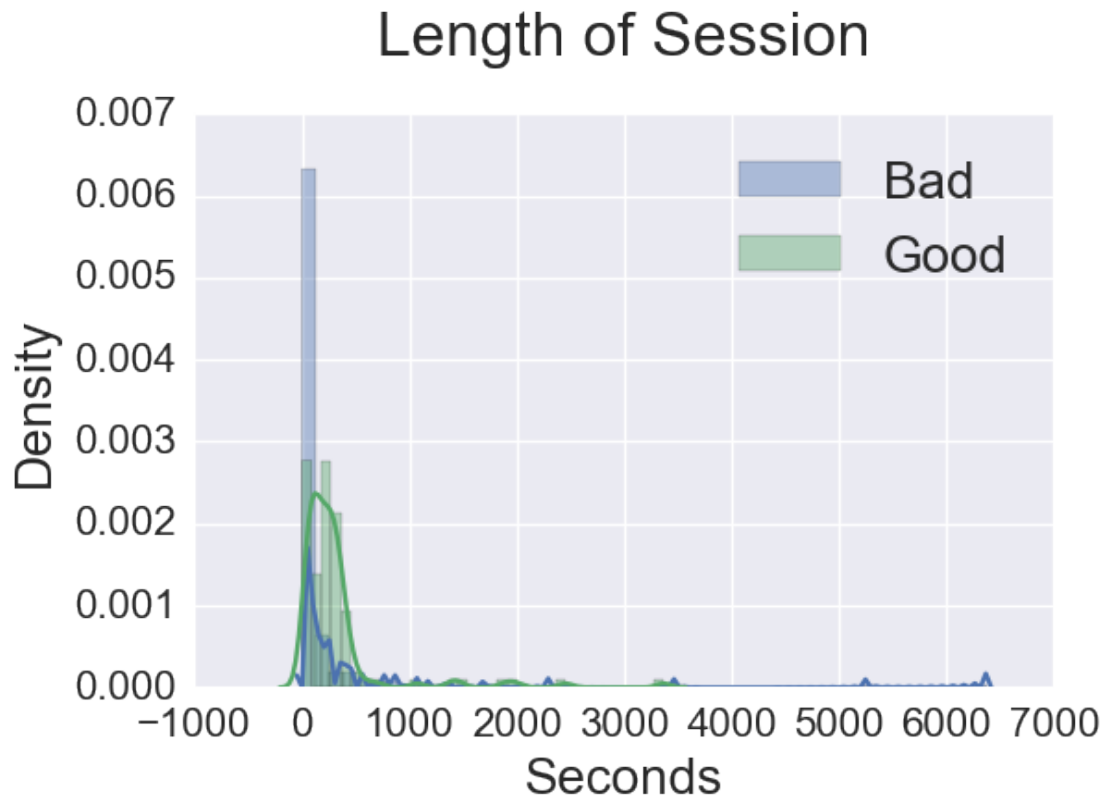


```
In [74]: sns.distplot( df4_t['avg_data_mean'],label = 'True')
          sns.distplot( df4_f['avg_data_mean'], label = 'False')
          plt.legend()
```

```
Out[74]: <matplotlib.legend.Legend at 0x12c8afd90>
```




```
In [75]: sns.distplot( df4_t['avg_data_len']/60.0, label = 'Bad')
sns.distplot( df4_f['avg_data_len']/60.0, label = 'Good')
plt.xlabel('Seconds', fontsize=20)
plt.ylabel('Density', fontsize=20)
plt.title('Length of Session', fontsize=24, y=1.08)
plt.tick_params(axis='both', labels=16)
plt.legend(fontsize=20)
plt.savefig('feat-session-length.png', bbox_inches='tight')
```




```
In [76]: sns.distplot( df4_t['avg_data_max'],label = 'Bad')
          sns.distplot( df4_f['avg_data_max'], label = 'Good')
          plt.legend(loc=0)
```

```
Out [76]: <matplotlib.legend.Legend at 0x1207f61d0>
```



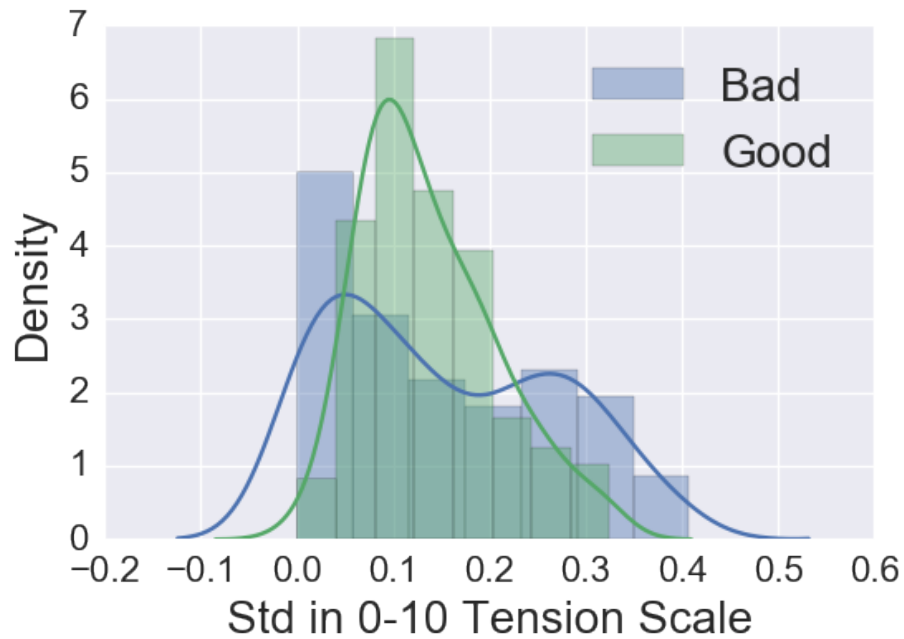
```
In [77]: sns.distplot( df4_t['avg_data_min'], label = 'Bad')  
          sns.distplot( df4_f['avg_data_min'], label = 'Good')  
          plt.legend()
```

```
Out[77]: <matplotlib.legend.Legend at 0x120a3efd0>
```



```
In [78]: sns.distplot( df4_t['avg_data_std']/4095.0, label = 'Bad')
sns.distplot( df4_f['avg_data_std']/4095.0, label = 'Good')
plt.xlabel('Std in 0-10 Tension Scale', fontsize=20)
plt.ylabel('Density', fontsize=20)
plt.title('Standard Deviation in Tension Measurement', fontsize=24, y=1.08)
plt.tick_params(axis='both', labelsize=16)
plt.legend(fontsize=20)
plt.savefig('feat-std.png', bbox_inches='tight')
```

Standard Deviation in Tension Measurement



```
In [79]: good = [len(np.where(df4_f['avg_data_std']<=60)[0]),
                  len(np.where(df4_f['avg_data_std']>60)[0]) ]
```

```
bad = [len(np.where(df4_t['avg_data_std']<=60)[0]),
        len(np.where(df4_t['avg_data_std']>60)[0]) ]
```

```
In [80]: fig, ax = plt.subplots()
```

```
index = np.arange(2)
bar_width = 0.35
```

```
opacity = 0.4
error_config = {'ecolor': '0.3'}
```

```
rects1 = plt.bar(index,
                  bad,
                  bar_width,
                  alpha=opacity,
                  color='b',
                  label='Bad')
```

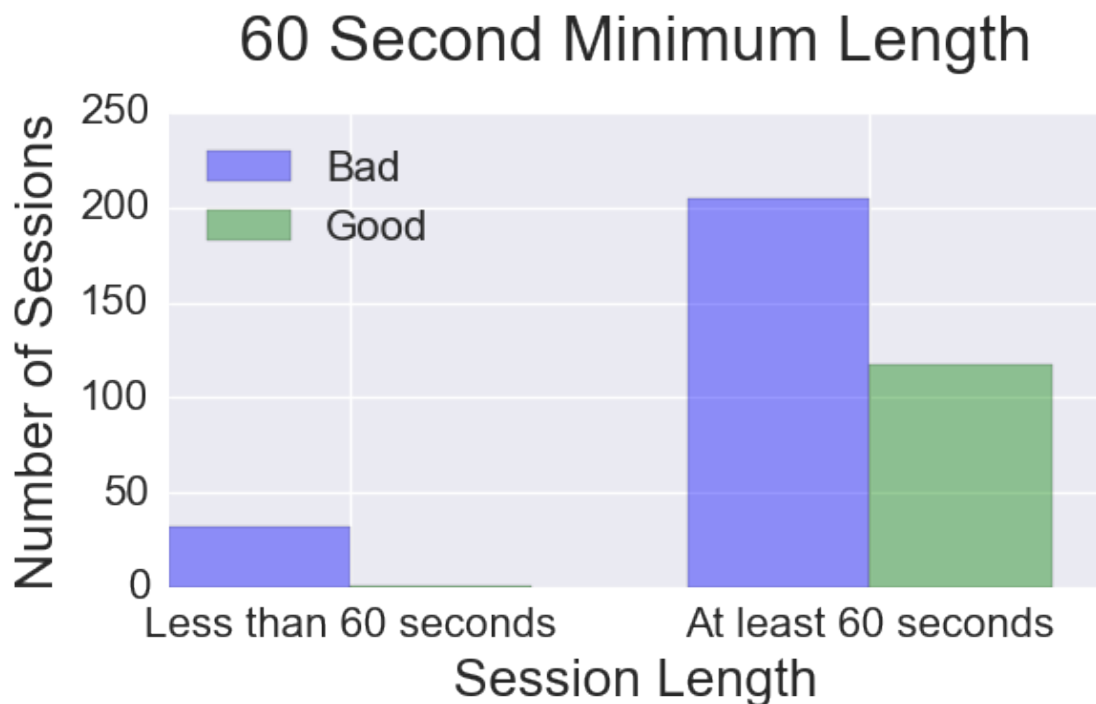
```
rects2 = plt.bar(index + bar_width,
```

```

        good,
        bar_width,
        alpha=opacity,
        color='g',
        label='Good')

plt.xlabel('Session Length', fontsize=20)
plt.ylabel('Number of Sessions', fontsize=20)
plt.title('60 Second Minimum Length', fontsize=24, y=1.08)
plt.xticks(index + bar_width, ('Less than 60 seconds', 'At least 60 seconds'))
plt.legend(fontsize=16, loc=0)
plt.tick_params(axis='both', labelsize=16)
plt.tight_layout()
plt.show()
plt.savefig('feat-60-sec_min.png', bbox_inches='tight')

```



<matplotlib.figure.Figure at 0x120f3a510>


```

In [81]: sns.distplot(df4_t['pegged_low'], kde=False, label = 'Bad')
         sns.distplot(df4_f['pegged_low'], kde=False, label = 'Good')

         plt.legend()


```

Out[81]: <matplotlib.legend.Legend at 0x1214410d0>



```
In [82]: sns.distplot( df4_t['pegged_high'],kde=False, label = 'Bad')  
sns.distplot( df4_f['pegged_high'], kde=False, label = 'Good')  
plt.legend()
```

```
Out[82]: <matplotlib.legend.Legend at 0x1224f7610>
```

```
In [83]: # Set up the matplotlib figure
df4.columns
```

```
Out[83]: Index([ u'Session_name',      u'userid',      u'score', u'avg_data_me
               u'avg_data_max', u'avg_data_min', u'avg_data_len', u'avg_data_s
               u'session_num',    u'min_len',    u'pegged_low',  u'pegged_hi
               u'cycles_1',      u'swings_1'],
              dtype='object')
```

```
In [84]: clf = RandomForestClassifier(n_estimators=10, random_state=42).fit(X_train

      scores = cross_val_score(clf, X_train, y_train, cv=10)
      print scores
      print np.average(scores)
```

```
[ 0.79166667  0.95833333  0.875      0.875      0.875      0.95833333
  0.91666667  0.95833333  0.82608696  0.82608696]
0.886050724638
```

```
In [85]: joblib.dump(clf, 'brtakrf_class.pkl')
```

```
Out[85]: ['brtakrf_class.pkl']
```

```
In [ ]:
```

```

In [86]: clfab = AdaBoostClassifier(base_estimator=None,
    n_estimators=50, learning_rate=1.0, algorithm='SAMME.R',
    random_state=42).fit(X_train,y_train)

In [87]: scores_ab = cross_val_score(clfab, X_train, y_train, cv=10)
    print scores_ab
    print np.average(scores_ab)

[ 0.79166667  0.91666667  0.95833333  0.875          0.875          0.875          0.875
  0.875          0.7826087   0.7826087 ]
0.860688405797

In [88]: clfgb = GradientBoostingClassifier(loss='deviance', learning_rate=0.1,
    n_estimators=100, subsample=1.0,
    criterion='friedman_mse', min_samples_split=2,
    min_samples_leaf=1, min_weight_fraction_leaf=0.0,
    max_depth=3, min_impurity_split=1e-07, init=None,
    random_state=42, max_features=None,
    verbose=0, max_leaf_nodes=None,
    warm_start=False, presort='auto').fit(X_train,y_train)

In [89]: scores_gb = cross_val_score(clfgb, X_train, y_train, cv=10)
    print scores_gb
    print np.average(scores_gb)

[ 0.875          0.91666667  0.91666667  0.875          0.91666667  0.91666667
  0.83333333  0.91666667  0.86956522  0.91304348]
0.894927536232

In [90]: scores_rf = cross_val_score(clf, X_train, y_train, cv=10)
    print scores_rf

[ 0.79166667  0.95833333  0.875          0.875          0.875          0.95833333
  0.91666667  0.95833333  0.82608696  0.82608696]

In [91]: RF_pred = clf.predict(X2)

In [92]: df4['RF_pred'] = RF_pred

In [93]: y2 = df4['score'].as_matrix()
    X_df2 = df4[['avg_data_mean',
    'avg_data_max', 'avg_data_min', 'avg_data_std', 'avg_data_len',
    'cycles_1','swings_1','min_len','pegged_low','pegged_high','RF
    X2 = X_df2.as_matrix()

In [94]: X_train, X_test, y_train, y_test = train_test_split(
    X2, y2, test_size=0.33, random_state=42)

```

```
In [95]: clfab = AdaBoostClassifier(base_estimator=None, n_estimators=50,
                                   learning_rate=1.0, algorithm='SAMME.R', random_state=42).fit(X_train, y_train)

scores_ab = cross_val_score(clfab, X_train, y_train, cv=10)
print scores_ab
print np.average(scores_ab)
```

1.	1.	1.	0.95833333	1.	1.
0.95833333	1.	1.	0.95652174]		

0.98731884058

```
In [96]: scores_ab = cross_val_score(clfab, X_test, y_test, cv=10)
print scores_ab
print np.average(scores_ab)
```

0.91666667	0.83333333	1.	0.91666667	1.	0.91666667
0.91666667	0.91666667	1.	0.9		

0.931666666667

```
In [97]: clfgb = GradientBoostingClassifier(loss='deviance',
                                             learning_rate=0.1, n_estimators=100, subsample=1.0,
                                             criterion='friedman_mse', min_samples_split=2,
                                             min_samples_leaf=1, min_weight_fraction_leaf=0.0,
                                             max_depth=3, min_impurity_split=1e-07, init=None,
                                             random_state=42, max_features=None, verbose=0,
                                             max_leaf_nodes=None, warm_start=False, presort='auto').fit(X_train, y_train)

scores_gb = cross_val_score(clfgb, X_train, y_train, cv=10)
print scores_gb
print np.average(scores_gb)
```

1.	1.	1.	0.95833333	1.	1.	1.
1.	1.	0.95652174]				

0.991485507246

```
In [98]: scores_gb = cross_val_score(clfgb, X_test, y_test, cv=10)
print scores_gb
print np.average(scores_gb)
```

0.83333333	0.83333333	0.91666667	1.	1.	0.91666667
0.91666667	1.	1.	0.9		

0.931666666667

```
In [99]: joblib.dump(clf, 'brtakgb_class.pkl')
```

```
Out[99]: ['brtakgb_class.pkl']
```

```

In [100]: trial_users = pd.read_csv('TrialUsers.csv', header=None)

In [101]: df_trial = df4[df_results['userid'].isin(trial_users[0]) | df_results['use

In [102]: y_trial = df_trial['score'].as_matrix()
          X_trial_df = df_trial[['avg_data_mean',
                                'avg_data_max', 'avg_data_min', 'avg_data_std', 'avg_data_len',
                                'cycles_1', 'swings_1', 'min_len', 'pegged_low', 'pegged_high']]

In [103]: X_trial = X_trial_df.as_matrix()

In [104]: trial_score = clf.score(X_trial, y_trial)
          print trial_score

0.96875

In [105]: X_trial_df2= X_trial_df.copy()
          X_trial_df2['RF_pred'] = clf.predict(X_trial)
          X_trial2 = X_trial_df2.as_matrix()
          X_train, X_test, y_train, y_test = train_test_split(
              X_trial2, y_trial, test_size=0.33, random_state=42)

In [106]: trial_score_gb = clfgb.score(X_trial2, y_trial)
          print trial_score_gb

0.96875

In [107]: trial_score_ab = clfab.score(X_trial2, y_trial)
          print trial_score_ab

0.96875

In [108]: clfab.feature_importances_

Out[108]: array([ 0.04,  0.14,  0.06,  0.22,  0.08,  0. ,  0.1 ,  0. ,  0. ,
                  0. ,  0.36])

In [109]: len(X_trial[0])

Out[109]: 10

In [110]: print df_trial.columns
          len(df_trial.columns)

Index([ u'Session_name',          u'userid',          u'score', u'avg_data_mean',
        u'avg_data_max',  u'avg_data_min',  u'avg_data_len',  u'avg_data_std',
        u'session_num',    u'min_len',    u'pegged_low',    u'pegged_high',
        u'cycles_1',      u'swings_1',      u'RF_pred'],
      dtype='object')

```

```
Out[110]: 15
```

```
In [111]: print X_df.columns
          print X_df2.columns
          print X_trial_df.columns
          print X_trial_df2.columns
```

```
Index([u'avg_data_mean', u'avg_data_max', u'avg_data_min', u'avg_data_len',
       u'avg_data_std', u'session_num', u'min_len', u'pegged_low',
       u'pegged_high'],
      dtype='object')
```

```
Index([u'avg_data_mean', u'avg_data_max', u'avg_data_min', u'avg_data_std',
       u'avg_data_len', u'cycles_1', u'swings_1', u'min_len', u'pegged_low',
       u'pegged_high', u'RF_pred'],
      dtype='object')
```

```
Index([u'avg_data_mean', u'avg_data_max', u'avg_data_min', u'avg_data_std',
       u'avg_data_len', u'cycles_1', u'swings_1', u'min_len', u'pegged_low',
       u'pegged_high'],
      dtype='object')
```

```
Index([u'avg_data_mean', u'avg_data_max', u'avg_data_min', u'avg_data_std',
       u'avg_data_len', u'cycles_1', u'swings_1', u'min_len', u'pegged_low',
       u'pegged_high', u'RF_pred'],
      dtype='object')
```

```
In [112]: print 'Bad=',len([x for x in y_trial if x==True])
          print 'Good=',len([x for x in y_trial if x==False])
```

```
Bad= 33
Good= 31
```

```
In [113]: y_pred2 = clfgb.predict(X_trial2)
```

```
In [114]: print 'Bad=',len([x for x in y_pred2 if x==True])
          print 'Good=',len([x for x in y_pred2 if x==False])
```

```
Bad= 31
Good= 33
```

```
In [115]: right = y_trial == y_pred2
```

```
In [116]: print 'Correct=',len([x for x in right if x==True])
          print 'Incorrect=',len([x for x in right if x==False])
```

```
Correct= 62
Incorrect= 2
```

```
In [117]: y_pred2 = clf.predict(X_trial)
          right = y_trial == y_pred2
          print 'Correct=',len([x for x in right if x==True])
          print 'Incorrect=',len([x for x in right if x==False])
```

```
Correct= 62
Incorrect= 2
```

```
In [118]: #df_trial[df_trial['score']!=df_trial['RF_pred']]
```

```
In [119]: # Compute the correlation matrix
          corr = df_results.corr()
```

```
          # Generate a mask for the upper triangle
          mask = np.zeros_like(corr, dtype=np.bool)
          mask[np.triu_indices_from(mask)] = True
```

```
          # Set up the matplotlib figure
          f, ax = plt.subplots(figsize=(11, 9))
```


```
          # Generate a custom diverging colormap
          cmap = sns.diverging_palette(220, 10, as_cmap=True)
```

```
          # Draw the heatmap with the mask and correct aspect ratio
```

```
          sns.heatmap(corr, mask=mask, cmap=cmap, vmax=.3,
                      square=True,
```

```
                      linewidths=.5, cbar_kws={"shrink": .5}, ax=ax, annot=True,
                      plt.xticks(rotation=45)
                      plt.yticks(rotation=0))
```

```
Out[119]: (array([ 0.5,  1.5,  2.5,  3.5,  4.5,  5.5,  6.5,  7.5,  8.5,  9.5]),
          <a list of 10 Text yticklabel objects>)
```



```
In [120]: # Compute the correlation matrix
corr = df4.corr()

# Generate a mask for the upper triangle
mask = np.zeros_like(corr, dtype=np.bool)
mask[np.triu_indices_from(mask)] = True


# Set up the matplotlib figure
f, ax = plt.subplots(figsize=(11, 9))

# Generate a custom diverging colormap
cmap = sns.diverging_palette(220, 10, as_cmap=True)

# Draw the heatmap with the mask and correct aspect ratio
sns.heatmap(corr, mask=mask, cmap=cmap, vmax=.3,
            square=True,
            linewidths=.5, cbar_kws={"shrink": .5}, ax=ax, annot=True, an
```

```
plt.xticks(rotation=45)
plt.yticks(rotation=0)
```

Out[120]: (array([0.5, 1.5, 2.5, 3.5, 4.5, 5.5, 6.5, 7.5, 8.5, 9.5, 10.5, 11.5, 12.5]), <a list of 13 Text yticklabel objects>)



```
In [121]: # Compute the correlation matrix
df9=df4.drop('RF_pred',axis=1)
corr = df9.corr()

# Generate a mask for the upper triangle
mask = np.zeros_like(corr, dtype=np.bool)
mask[np.triu_indices_from(mask)] = True

# Set up the matplotlib figure
f, ax = plt.subplots(figsize=(11, 9))
```



```

# Generate a custom diverging colormap
cmap = sns.diverging_palette(220, 10, as_cmap=True)


# Draw the heatmap with the mask and correct aspect ratio
sns.heatmap(corr, mask=mask, cmap=cmap, vmax=.3,
            square=True,
            linewidths=.5, cbar_kws={"shrink": .5}, ax=ax, annot=True, an
plt.xticks(rotation=45)
plt.yticks(rotation=0)

```

```

Out[121]: (array([ 0.5,  1.5,  2.5,  3.5,  4.5,  5.5,  6.5,  7.5,  8.5,
                    9.5, 10.5, 11.5]), <a list of 12 Text yticklabel objects>)

```



```

In [122]: forest = clf_save
          names = clf_name_save
          X = X_save

```

```

importances = forest.feature_importances_
print len(importances), len(forest.feature_importances_)
std = np.std([tree.feature_importances_ for tree in forest.estimators_],
              axis=0)

indices = np.argsort(importances)[::-1]

# Print the feature ranking
print("Feature ranking:")

for f in range(X.shape[1]):
    print("%d. feature %s (%f)" % (f + 1, names[indices[f]], importances[indices[f]]))

# Plot the feature importances of the forest
plt.figure()
plt.title("Feature importances")
plt.bar(range(X.shape[1]), importances[indices],
        color="r", yerr=std[indices], align="center")
plt.xticks(range(X.shape[1]), names[indices], rotation=45)
plt.xlim([-1, X.shape[1]])
plt.show()
plt.savefig('feat-imp.png',bbox_inches='tight')

```

10 10

Feature ranking:

1. feature avg_data_len (0.251626)
2. feature cycles_1 (0.134859)
3. feature min_len (0.134769)
4. feature avg_data_std (0.105646)
5. feature avg_data_min (0.097373)
6. feature pegged_low (0.094684)
7. feature avg_data_mean (0.070648)
8. feature swings_1 (0.057275)
9. feature avg_data_max (0.050543)
10. feature pegged_high (0.002578)



<matplotlib.figure.Figure at 0x123a8e690>